

```
# We Import two libraries(one standard and one third-party) needed for  
performing the linear algebraic manipulations  
# and the random initialization of the main variables(weights and  
biases). Finally, we make  
# use of the matplotlib library for monitoring the performance of the  
model and in turn selecting the adequate parameters.
```

```
import numpy as np  
import random  
import matplotlib.pyplot as plt
```

```
# The folowing neural network algorithm is built for classifying hand-  
written digits extracted from the widely known  
# MNIST dataset. This dataset contains a training, validation and test  
data totally comprised of 80.000 handwritten  
# digits that we aim to classify using Stochastic Gradient Descent and  
Backpropagation for making the changes in the  
# parameters as controllable as possible and the sigmoid function for  
deciding the activation or not of a neuron.  
# Lots of details are given along the construction of the model below.
```

```
class Network():
```

```
# Here we initialize the weights between the input layer and the  
first hidden layer of the network.
```

```
# Also we initialize the biases that the neurons in the first  
hidden layer of the network have. Note that
```

```
# the input layer has no biases.
```

```
# Finally sizes is a list indicating the number of neurons each  
layer has and so the length of the list
```

```
# determines the number of network's layers.
```

```
def __init__(self, sizes):  
    self.n_max_layers = len(sizes)  
    self.sizes = sizes  
    self.biases = [np.random.randn(x,1) for x in sizes[1:]]  
    self.weights = [np.random.randn(y,x) for x,y in zip(sizes[:-  
1],sizes[1:])] 
```

```
# To take the result of a single neuron based on sigmoid function  
(To get one prediction for the input a).
```

```
# Observe that this function updates the input value a as many  
times as the layers that we have, since
```

```
# zip(self.biases,self.weights) contains one tuple for each layer,  
having the biases as the first component
```

```
# and the respective weights from the previous to the following  
layer.
```

```
def feed_forward(self, a):  
    for b, w in zip(self.biases, self.weights):  
        a = sigmoid(np.dot(w,a)+b)  
    return a
```

Stochastic Gradient Descent(the process is using backprop function(Backpropagation) which performs the gradient finding and will be explained afterwards). We use Stochastic gradient descent and not typical gradient decent in order to balance the trade-off between the large training data and the time needed to train the model.Namely, we partition the (shuffled)training data in mini batches and then we train the algorithm on each one of them separately.

In terms of the variables, epochs variable determines how many times the training data will be partioned into mini batches.

Moreover, eta is the learning rate used for the gradient descent where here is performed by update_mini_batch

which gives the updated position of the the weights and biases towards the minimization of the

cost function under consideration(the Quadratic Cost on this occasion).

```
def SGD(self,training_data, mini_batch_size, epochs, eta,
test_data = None):
```

```
    if test_data: n_test = len(test_data)
    n = len(training_data)
    for j in range(epochs):
        random.shuffle(training_data)
        mini_batches = [
            training_data[k:k+mini_batch_size]
            for k in range(0,n,mini_batch_size)]
        for mini_batch in mini_batches:
            self.update_mini_batch(mini_batch, eta)
        if test_data:
            print ('Epoch{0}: {1}/{2}'.format(
                j,self.evaluate(test_data),n_test))
        else:
            print ('Epoch{0}: Complete'.format(j))
```

The following function serves as a gradient descent iterator for all the training data contained in a mini-batch.

Actually what it does is that it first collects all the partial derivatives of the cost function with respect to

biases and weights for all the train data inside the mini batch using backpropagation and finally updates the biases

and weights by performing a gradient descent based on the given learning rate(eta).

```
def update_mini_batch(self,mini_batch,eta):
    gradient_w = [np.zeros(w.shape) for w in self.weights]
    gradient_b = [np.zeros(b.shape) for b in self.biases]
    for x,y in mini_batch:
        delta_gd_w , delta_gd_b = self.backprop(x,y)
        gradient_w = [gd_w+dgd_w for gd_w, dgd_w in
zip(gradient_w, delta_gd_w)]
```

```

        gradient_b = [gd_b+dgd_b for gd_b, dgd_b in
zip(gradient_b, delta_gd_b)]
        self.weights = [w - (eta/len(mini_batch))*gd_w for w, gd_w in
zip(self.weights, gradient_w)]
        self.biases = [b - (eta/len(mini_batch))*gd_b for b, gd_b in
zip(self.biases, gradient_b)]

# Backpropagation is method allowing us to determine the partial
derivatives of the cost function with respect to
# weights and biases. Here, the function takes on a single
training data point (x,y) as an input.
def backprop(self,x,y):
    gradient_b = [np.zeros(b.shape) for b in self.biases]#
Collection of Partial Derivative of the cost function with respect to
b(bias)
    gradient_w = [np.zeros(w.shape) for w in self.weights]#
Collection of Partial Derivative of the cost function with respect to
w(weights)
    activation = x #Input is by default the first activation value
    activations = [x] # Here we will collect all the activations
    zs = [] # To collect the z_values given in turn to the sigmoid
function as input.
    for b,w in zip(self.biases, self.weights):
        z = np.dot(w,activation) + b
        zs.append(z)
        activation = sigmoid(z)
        activations.append(activation)
    delta = (activations[-1] - y) * sigmoid_prime(zs[-1])
    gradient_b[-1] = delta
    gradient_w[-1] = np.dot(delta, activations[-2].transpose())
    for l in range(2, self.n_max_layers):
        delta = np.dot(self.weights[-l+1].transpose(),gradient_b[-
l+1]) * sigmoid_prime(zs[-l])
        gradient_b[-l] = delta
        gradient_w[-l] = np.dot(delta,activations[-l-
1].transpose())
    return (gradient_w, gradient_b)

# A fuction calculating the correct model predictions on the train
data, if provided. The predictions are based on the
# highest value given by the sigmoid function between the neurons
in the output layer.
def evaluate(self,test_data):
    test_results = [(np.argmax(self.feed_forward(x)),y)
                     for x,y in test_data]
    return sum(int(x==y) for x,y in test_results)

# The sigmoid function used for the decision making of the neurons

```

*taking as input a variable of the form $z = \Sigma w * x + b$, where w represents the weights, x the input values and b the bias.*

```
def sigmoid(z):  
    return 1.0/(1.0 + np.exp(-z))  
  
# The Derivative of the sigmoid function  
def sigmoid_prime(z):  
    return (sigmoid(z) * (1 - sigmoid(z)))
```

The following three functions are used for loading the dataset under investigation(MNIST) in an appropriate format.

*# Firstly we load the dataset using the load_data function and then we create the load_data_wrapper for mainly
turn all inputs into (784,1) matrices since the inputs are 28x28(pixels)images. Moreover, in the training data we turn
the actual hand-written digits(y) into a binary vector*

```
import gzip  
import pickle
```

```
def load_data():  
    f = gzip.open('mnist.pkl.gz', 'rb')  
    u = pickle._Unpickler( f )  
    u.encoding = 'latin1'  
    train, val, test = u.load()  
    return (train, val, test)
```

```
def load_data_wrapper():  
    tr,val,test = load_data()  
    tr_1 = [x.reshape(784,1) for x in tr[0]]  
    tr_2 = [vectorize(y) for y in tr[1]]  
    training_data = zip(tr_1, tr_2)  
    training_data_fixed = [x for x in training_data]  
    val_1 = [x.reshape(784,1) for x in val[0]]  
    validation_data = zip(val_1, val[1])  
    validation_data_fixed = [x for x in validation_data]  
    tst_1 = [x.reshape(784,1) for x in test[0]]  
    test_data = zip(tst_1, test[1])  
    test_data_fixed = [x for x in test_data]  
    return (training_data_fixed, validation_data_fixed,  
test_data_fixed)
```

```
def vectorize(y):  
    e = np.zeros((10,1))  
    e[y] = 1.0  
    return e
```

```
training_data, validation_data, test_data = load_data_wrapper()
```

*# Now, we are ready to perform the model and take an account of its classification success rate for different parameters.
Note that, the neural network begins with 784 neurons in the input layer and 10 neurons in the ouput layer to be in line
with the (modified)structure of the dataset.*

```
net = Network([784,30,10])
```

```
net.SGD(training_data, mini_batch_size = 15, epochs = 30, eta = 0.5,  
test_data = validation_data)
```

```
Epoch0: 6973/10000  
Epoch1: 8470/10000  
Epoch2: 8774/10000  
Epoch3: 8930/10000  
Epoch4: 9016/10000  
Epoch5: 9072/10000  
Epoch6: 9108/10000  
Epoch7: 9131/10000  
Epoch8: 9174/10000  
Epoch9: 9179/10000  
Epoch10: 9212/10000  
Epoch11: 9228/10000  
Epoch12: 9245/10000  
Epoch13: 9245/10000  
Epoch14: 9268/10000  
Epoch15: 9281/10000  
Epoch16: 9280/10000  
Epoch17: 9279/10000  
Epoch18: 9300/10000  
Epoch19: 9315/10000  
Epoch20: 9318/10000  
Epoch21: 9322/10000  
Epoch22: 9314/10000  
Epoch23: 9335/10000  
Epoch24: 9344/10000  
Epoch25: 9354/10000  
Epoch26: 9342/10000  
Epoch27: 9356/10000  
Epoch28: 9353/10000  
Epoch29: 9367/10000
```

Now, we introduce the following three modifications to the built neural net above aiming to ameliorate its performance.

*# Backpropagation with the Cross_Entropy as the cost function. Cross-Entropy function allow the network to overcome
learning slowdowns. In practice, the rate of change of the cost function is not dependent on the derivative of the
sigmoid function which is responsible for the learning slowdown when*

the input values are close to either 0 or 1.
 # Lets state here the definition of the cross-entropy function: $C(a) = -\frac{1}{n} \sum_{x} (y \ln(a) + (1-y) \ln(1-a))$, where $y(=y(x))$ is
 # the actual value, $a(x)$ represents the outcome of the sigmoid
 function and the sum is taken over the training inputs x
 # which are n in total.

L1 and L2 Regulizations for ameliorating classification accuracy and
 reducing overfitting in neural nets. In both occasions
 # we add an extra component in the cost function which tries to keep
 weights magnitude low, or in other words to make the total
 # net less dependent on individual piece of evidence. Another
 Regulatory technique for neural nets is the Drop-Out Method. For
 # now we are going to focus on L2 and L1 regularizations and in the
 following we try to include them inside a new neural net.
 # Also, we will incorporate the cost function used in the model as a
 class variable with two methods(quadratic cost and cross
 # entropy).

Finally we take care of the weights and biases initialization.
 Namely, we try to keep the distribution of the the variable
 # $\sum w \cdot x + b$, where w are the weights, x the input values and b the
 biases, as close to $\text{Gaussian}(0,1)$ as possible. Loosely speaking,
 # if the input values were either 0 or 1 then we only need to control
 the variance of the weights and to do so we divide
 # we divide each weight variable with the square root of the amount of
 weights.

```
class initialize():
```

```
    def random_initialize(sizes):
        biases = [np.random.randn(x,1) for x in sizes[1:]]
        weights = [np.random.randn(y,x) for x,y in zip(sizes[:-1],sizes[1:])]
        return (biases, weights)
```

```
    def standarised_initialize(sizes):
        biases = [np.random.randn(x,1) for x in sizes[1:]]
        weights = [np.random.randn(y,x)/np.sqrt(x) for x,y in
zip(sizes[:-1],sizes[1:])]
        return (biases,weights)
```

```
class QuadraticCost():
```

```
    def fun(a,y):
        return (1/2)*np.linalg.norm(y-a)^2
```

```
    def delta(z,a,y):
        return (a-y)*sigmoid_prime(a)
```

```

class cross_entropy():

    def fun(a,y):
        return np.sum(np.nan_to_num(-y*np.log(a) - (1-y)*np.log(1-a)))

    def delta(z,a,y):
        return (a-y)

```

In the following neural net, except from adding new techniques that were described above, we create flags in the
Stochastic Gradient Descent function, indicating which metric results for the model we want to generate and in turn
enabling the plotting of these results for many different model (hyper-)parameters as shown in the following steps.

```

class Network1():

    # Here we initialize the weights between the input layer and the first hidden layer of the network.
    # Also we initialize the biases that the neurons in the first hidden layer of the network have. Note that
    # the input layer has no biases, and thats why start with 1 to be the first index used from sizes.
    # Finally sizes is a list indicating the number of neurons each layer has and in turn the length of the list
    # determines the number of network's layers.
    def __init__(self,sizes,cost = cross_entropy):
        self.n_max_layers = len(sizes)
        self.sizes = sizes
        self.cost = cost
        self.biases, self.weights =
initialize.standarised_initialize(sizes)

    # To take the result of a single neuron based on sigmoid function (To get one prediction for the input a)
    # Observe that this function updates the input value (a) as many times as the layers that we have, since
    # zip(self.biases,self.weights) contains one tuple for each layer, having the biases as the first component
    # and the respective weights from the previous layer to the following one.
    def feed_forward(self, a):
        for b, w in zip(self.biases, self.weights):
            a = sigmoid(np.dot(w,a)+b)
        return a

```

```

# Stochastic Gradient Descent(the process is using backprop
function(Backpropagation) which performs the gradient
# finding and will be explained afterwards). We use Stochastic
gradient descent in order to balance the trade-off
# between the large training data and the time needed to train the
model.Namely, we partition the (shuffled)training
# data in mini batches and then we train the algorithm on each one
of them separately.
# In terms of the variables, epochs variable determines how many
times the training data will be partioned into mini batches.
# Moreover, eta is the learning rate used for the gradient descent
where here is performed by update_mini_batch
# which gives the updated position of the input(based on the
weights and biases) towards the minimization of the
# cost function under consideration(MSE on this occasion).
def SGD(self,training_data, mini_batch_size, epochs, eta,lmbd,
    monitor_training_cost = False,
    monitor_training_accuracy = False,
    monitor_evaluation_cost = False,
    monitor_evaluation_accuracy = False,
    plot_accuracy = False,
    plot_cost = False,
    hyper_training_cost_plot = False,# The term 'hyper' is
used to indicate that these variables are for plots for
hyperparameters determination
    hyper_training_accuracy_plot = False,
    hyper_evaluation_cost_plot = False,
    hyper_evaluation_accuracy_plot = False,
    evaluation_data = None):
    training_cost = []
    training_accuracy = []
    evaluation_cost = []
    evaluation_accuracy = []
    if evaluation_data: n_evaluation = len(evaluation_data)
    n = len(training_data)
    for j in range(epochs):
        random.shuffle(training_data)
        mini_batches = [
            training_data[k:k+mini_batch_size]
            for k in range(0,n,mini_batch_size)]
        for mini_batch in mini_batches:
            self.update_mini_batch(mini_batch, eta, lmbd,n)
        if monitor_training_cost:
            cost = self.total_cost(training_data,lmbd, convert =
False)
            training_cost.append(cost)
            #print ("Cost on training data-Epoch{}: {} ".format(j,
                #cost))
        if monitor_training_accuracy:
            accuracy = self.accuracy(training_data, convert =

```



```

True)
        training_accuracy.append(accuracy)
        #print ("Accuracy on training data-Epoch{:}: {} /
{:}".format(j,
            #accuracy, n))
        if monitor_evaluation_cost:
            cost = self.total_cost(evaluation_data, lmbd, convert =
False)
            evaluation_cost.append(cost)
            #print ("Cost on evaluation data-Epoch{:}: {}
{:}.format(j,
            #cost))
        if monitor_evaluation_accuracy:
            accuracy = self.accuracy(evaluation_data, convert =
False)
            evaluation_accuracy.append(accuracy)
            print ("Accuracy on evaluation data-Epoch{:}: {} /
{:}".format(j,
            accuracy, n_evaluation))
        .....,
        if evaluation_data:
            n_evaluation = len(evaluation_data)
            print ('Epoch{:}: {1}/{2}'.format(
                j, self.evaluate(evaluation_data), n_evaluation))
        else:
            print ('Epoch{:}: Complete'.format(j))
        .....,
        if plot_accuracy:
            fig, ax = plt.subplots(figsize = (12, 8))
            ax.plot([int(x) for x in range(epochs)], [x/n_evaluation
for x in evaluation_accuracy],
                    label = 'Evaluation Accuracy', color = 'red')
            ax.plot([int(x) for x in range(epochs)], [x/n for x in
training_accuracy],
                    label = 'Training Accuracy', color = 'yellow')
            ax.legend(loc = 0)
            ax.set_xlabel('Epochs')
            ax.set_ylabel('Accuracy')
            plt.xticks(list(range(epochs)), range(epochs))
            plt.plot()
        if plot_cost:
            fig, ax = plt.subplots(figsize = (12, 9))
            fig1, ax1 = plt.subplots(figsize = (12, 9))
            ax.plot([int(x) for x in range(epochs)], [x for x in
evaluation_cost], label = 'Evaluation Cost', color = 'red')
            ax1.plot([int(x) for x in range(epochs)], [x for x in
training_cost], label = 'Training Cost', color = 'yellow')
            ax.legend(loc = 0)
            ax.set_xlabel('Epochs', weight = 'bold')
            ax.set_ylabel('Cost', weight = 'bold')

```

```

        ax.set_title('Validation data', weight = 'bold')
        ax1.legend(loc = 0)
        ax1.set_xlabel('Epochs', weight = 'bold')
        ax1.set_ylabel('Cost', weight = 'bold')
        ax1.set_title('Training data', weight = 'bold')
        plt.plot()
    if hyper_training_cost_plot:
        return training_cost
    if hyper_training_accuracy_plot:
        return training_accuracy
    if hyper_evaluation_accuracy_plot:
        return evaluation_accuracy
    if hyper_evaluation_cost_plot:
        return evaluation_cost

# The difference in updating the weights and the biases to the
previous model is that now we also
# add the regularization term 'lmbd'.
    def update_mini_batch(self, mini_batch, eta, lmbd, n):
        gradient_w = [np.zeros(w.shape) for w in self.weights]
        gradient_b = [np.zeros(b.shape) for b in self.biases]
        for x, y in mini_batch:
            delta_gd_w, delta_gd_b = self.backprop(x, y)
            gradient_w = [gd_w + dgd_w for gd_w, dgd_w in
zip(gradient_w, delta_gd_w)]
            gradient_b = [gd_b + dgd_b for gd_b, dgd_b in
zip(gradient_b, delta_gd_b)]
            # Here is where L2-Regularization comes into play
            #self.weights = [w - (eta/len(mini_batch))*gd_w -
(eta*lmbd*np.sign(w))/n for w, gd_w in zip(self.weights, gradient_w)]#
L1 Regularization
            self.weights = [w - (eta/len(mini_batch))*gd_w -
(eta*lmbd*w)/n for w, gd_w in zip(self.weights, gradient_w)]# L2
Regularization
            self.biases = [b - (eta/len(mini_batch))*gd_b for b, gd_b in
zip(self.biases, gradient_b)]

    # Again the backpropagation method with adjusted delta value
based on whether the cost function is the Quadratic or the
# Cross-Entropy.
    def backprop(self, x, y):
        gradient_b = [np.zeros(b.shape) for b in self.biases]#
Collection of Partial Derivative of the cost function with respect to
b(bias)
        gradient_w = [np.zeros(w.shape) for w in self.weights]#
Collection of Partial Derivative of the cost function with respect to
w(weights)
        activation = x #Input is by default the first activation value

```

```

        activations = [x]# Here we will collect all the neuron
activation values.
        zs = []# To collect the z_values given in turn to the sigmoid
function as input.
        for b,w in zip(self.biases, self.weights):
            z = np.dot(w,activation) + b
            zs.append(z)
            activation = sigmoid(z)
            activations.append(activation)
        delta = self.cost.delta(z,activations[-1],y)
        gradient_b[-1] = delta
        gradient_w[-1] = np.dot(delta, activations[-2].transpose())
        for l in range(2, self.n_max_layers):
            delta = np.dot(self.weights[-l+1].transpose(),gradient_b[-
l+1]) * sigmoid_prime(zs[-l])
            gradient_b[-l] = delta
            gradient_w[-l] = np.dot(delta,activations[-l-
1].transpose())
        return (gradient_w, gradient_b)

```

Nothing changes in the way the model evaluates itself.

```

def evaluate(self,test_data):
    test_results = [(np.argmax(self.feed_forward(x)),y)
                     for x,y in test_data]
    return sum(int(x==y) for x,y in test_results)

```

This is a new a function calculating the total cost of the model for a given dataset. Note that at the end we add the regulaziation terms.

```

def total_cost(self,data,lmbd,convert):
    cost = 0
    for x,y in data:
        a = self.feed_forward(x)
        if convert: y = vectorize(y)
        cost += self.cost.fun(a,y)/len(data)
    #cost += (lmbd/len(data))*sum(np.linalg.norm(w) for w in
self.weights)# L1 Regularization
    cost += 0.5*(lmbd/len(data))*sum(np.linalg.norm(w)**2 for w in
self.weights) # L2 Regularization
    return cost

```

An expansion of the evaluation function for performing two results collections based on whether is the training or validation(or test) set to be used. .

```

def accuracy(self,data,convert):
    results = []
    for x,y in data:

```

```

        a = self.feed_forward(x)
        if convert:
            results.append((np.argmax(a), np.argmax(y)))
        else:
            results.append((np.argmax(a), y))
    return sum(int(x == y) for x, y in results)

def sigmoid(z):
    return 1.0 / (1.0 + np.exp(-z))

def sigmoid_prime(z):
    return (sigmoid(z) * (1 - sigmoid(z)))

# Soft-max function instead of the sigmoid function for neurons
decision making. The attractive part of this function
# is that the sum of its outcome for all input values equals one.
Therefore, it can be seen as the Probability density function
# for the neurons in the output layer. Namely, we could think of it
as counting the probability that our estimation based on
# the input value belongs to one of the neurons (or activates the
neuron). Then, the above observation, affirms that the
# probability of the any neurons activation is 1. If we also employ
the negative log of the softmax function as a cost
# function, then the backpropagation delta parameter is identical to
the one obtained by using the cross-entropy as a cost
# function, as above.

def softmax(z):
    return np.exp(z) / sum(np.exp(z))

def log_cost(a):
    return -np.log(a)

array([-0.          , -0.          , -0.69314718, 75.98530807])

# Let us now run the model and check its accuracy on the training set.
The model we build here has two
# differences to the initial neural net above. In fact, the cost
function is taken to be the cross entropy and not the
# Quadratic cost and the initialization of weights and biases is
standarized and not random as explained earlier.
net = Network1(sizes = [784, 30, 10], cost = cross_entropy)
net.SGD(training_data, mini_batch_size = 15, epochs = 30, eta = 0.5,
        lmbd = 0,
        monitor_training_cost = False,

```

```
monitor_training_accuracy = True,  
monitor_evaluation_cost = False,  
monitor_evaluation_accuracy = False,  
evaluation_data = validation_data)
```

```
Accuracy on training data-Epoch0: 46472 / 50000  
Accuracy on training data-Epoch1: 47529 / 50000  
Accuracy on training data-Epoch2: 47978 / 50000  
Accuracy on training data-Epoch3: 48192 / 50000  
Accuracy on training data-Epoch4: 48205 / 50000  
Accuracy on training data-Epoch5: 48385 / 50000  
Accuracy on training data-Epoch6: 48258 / 50000  
Accuracy on training data-Epoch7: 48502 / 50000  
Accuracy on training data-Epoch8: 48427 / 50000  
Accuracy on training data-Epoch9: 48607 / 50000  
Accuracy on training data-Epoch10: 48740 / 50000  
Accuracy on training data-Epoch11: 48721 / 50000  
Accuracy on training data-Epoch12: 48812 / 50000  
Accuracy on training data-Epoch13: 48883 / 50000  
Accuracy on training data-Epoch14: 48890 / 50000  
Accuracy on training data-Epoch15: 48934 / 50000  
Accuracy on training data-Epoch16: 48887 / 50000  
Accuracy on training data-Epoch17: 49058 / 50000  
Accuracy on training data-Epoch18: 48979 / 50000  
Accuracy on training data-Epoch19: 49000 / 50000  
Accuracy on training data-Epoch20: 48885 / 50000  
Accuracy on training data-Epoch21: 49007 / 50000  
Accuracy on training data-Epoch22: 49232 / 50000  
Accuracy on training data-Epoch23: 49159 / 50000  
Accuracy on training data-Epoch24: 49219 / 50000  
Accuracy on training data-Epoch25: 49254 / 50000  
Accuracy on training data-Epoch26: 49227 / 50000  
Accuracy on training data-Epoch27: 49108 / 50000  
Accuracy on training data-Epoch28: 49200 / 50000  
Accuracy on training data-Epoch29: 49289 / 50000
```

We do the same but now we measure the accuracy on the validation set.

We can see that the Classification Success Rate becomes significantly better than it was with the Quadratic Cost

(from 9367/1000 to 9667/1000). However we can see that the model now reached high percentage accuracy from a very early on, and

from epoch 5 onwards the improvement was not really significant.

Also, compared to the training set above here there is no gradual increase and the results are fluctuating around 96%.

These facts imply possible overfitting of the neural net that we will try to tackle by an L2 regularization term in the

cost function.

In the following steps, we generate plots of model accuracy and cost

function on the training and validation set.

```
net = Network1(sizes = [784,30,10], cost = cross_entropy)
net.SGD(training_data, mini_batch_size = 15, epochs = 30, eta = 0.5,
        lmbd = 0,
            monitor_training_cost = False,
            monitor_training_accuracy = False,
            monitor_evaluation_cost = False,
            monitor_evaluation_accuracy = True,
            evaluation_data = validation_data)
```

```
Accuracy on evaluation data: 9387 / 10000
Accuracy on evaluation data: 9514 / 10000
Accuracy on evaluation data: 9511 / 10000
Accuracy on evaluation data: 9594 / 10000
Accuracy on evaluation data: 9622 / 10000
Accuracy on evaluation data: 9622 / 10000
Accuracy on evaluation data: 9643 / 10000
Accuracy on evaluation data: 9624 / 10000
Accuracy on evaluation data: 9514 / 10000
Accuracy on evaluation data: 9636 / 10000
Accuracy on evaluation data: 9654 / 10000
Accuracy on evaluation data: 9634 / 10000
Accuracy on evaluation data: 9645 / 10000
Accuracy on evaluation data: 9648 / 10000
Accuracy on evaluation data: 9649 / 10000
Accuracy on evaluation data: 9639 / 10000
Accuracy on evaluation data: 9659 / 10000
Accuracy on evaluation data: 9661 / 10000
Accuracy on evaluation data: 9662 / 10000
Accuracy on evaluation data: 9631 / 10000
Accuracy on evaluation data: 9633 / 10000
Accuracy on evaluation data: 9618 / 10000
Accuracy on evaluation data: 9667 / 10000
Accuracy on evaluation data: 9665 / 10000
Accuracy on evaluation data: 9663 / 10000
Accuracy on evaluation data: 9658 / 10000
Accuracy on evaluation data: 9653 / 10000
Accuracy on evaluation data: 9648 / 10000
Accuracy on evaluation data: 9661 / 10000
Accuracy on evaluation data: 9633 / 10000
```


Below is the plot of the above calculations, which shows that the discrepancy between the training and validation accuracy tend to increase as the epochs are passing.

```
net = Network1(sizes = [784,30,10], cost = cross_entropy)
net.SGD(training_data, mini_batch_size = 15, epochs = 30, eta = 0.5,
        lmbd = 0,
            monitor_training_cost = True,
            monitor_training_accuracy = True,
```

```

monitor_evaluation_cost = False,
monitor_evaluation_accuracy = True,
plot_accuracy = True,
evaluation_data = validation_data)

```



We also plot the progress of cost function from epoch to epoch both on training and validation data. From the cost plot on
training data, we can see that there is a decreasing trend generally which is connected with the almost constant amelioration
on the accuracy of the model on the training data which was seen earlier. Also observe that there are more fluctuations towards
the last epochs where the accuracy was not significantly improved (epoch 22: 49232 / 50000 and epoch 29: 49289 / 50000). This
an indication of overfitting that we will try to tackle by applying L2 Regularization. Regarding the cost on the validation set
there is an increasing trend which might explain the models' learning slowdown(Epoch 5 : 9622/10000, Epoch 29: 9633). But before going into regularization, lets try to experiment
with more learning rate values, taking as a comparison point the
one that we already use(i.e $\eta = 0.5$). Heuristically, we are trying out values larger and smaller than 0.5 by factor 10.

```

net = Network1(sizes = [784,30,10], cost = cross_entropy)
net.SGD(training_data, mini_batch_size = 15, epochs = 30, eta = 0.5,
lmbd = 0,

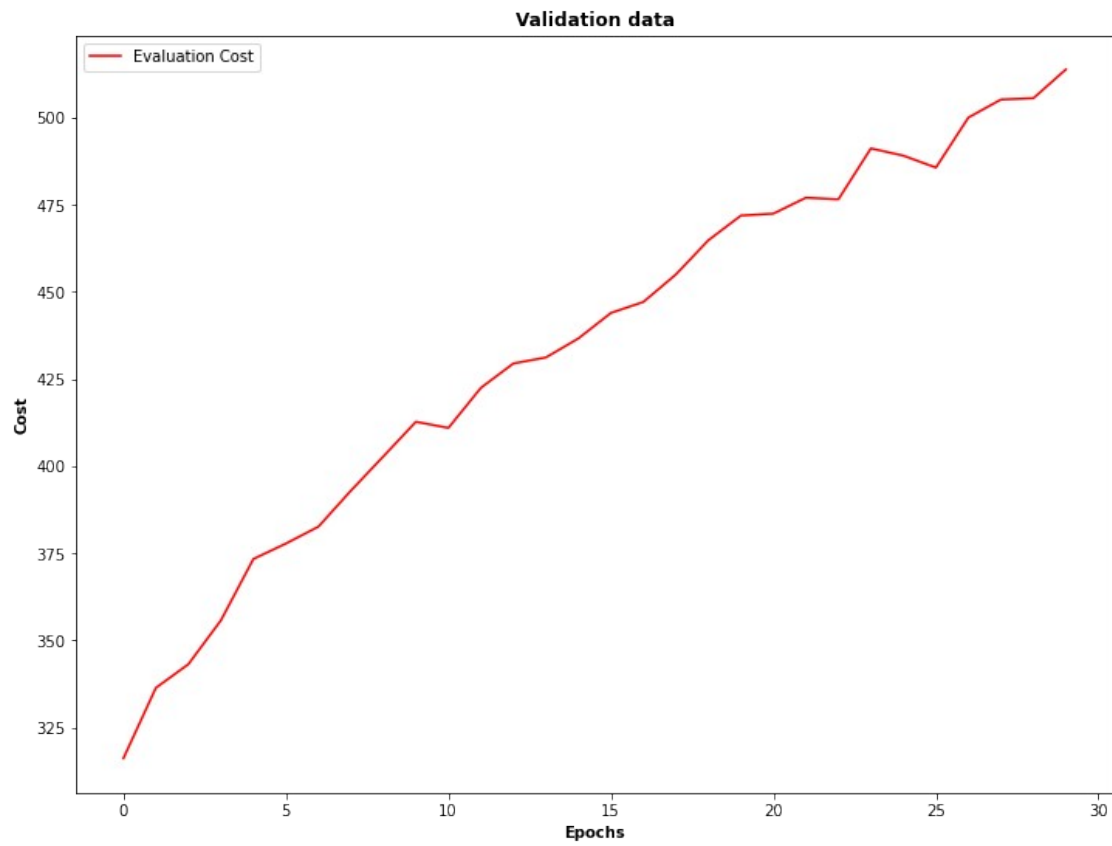
```

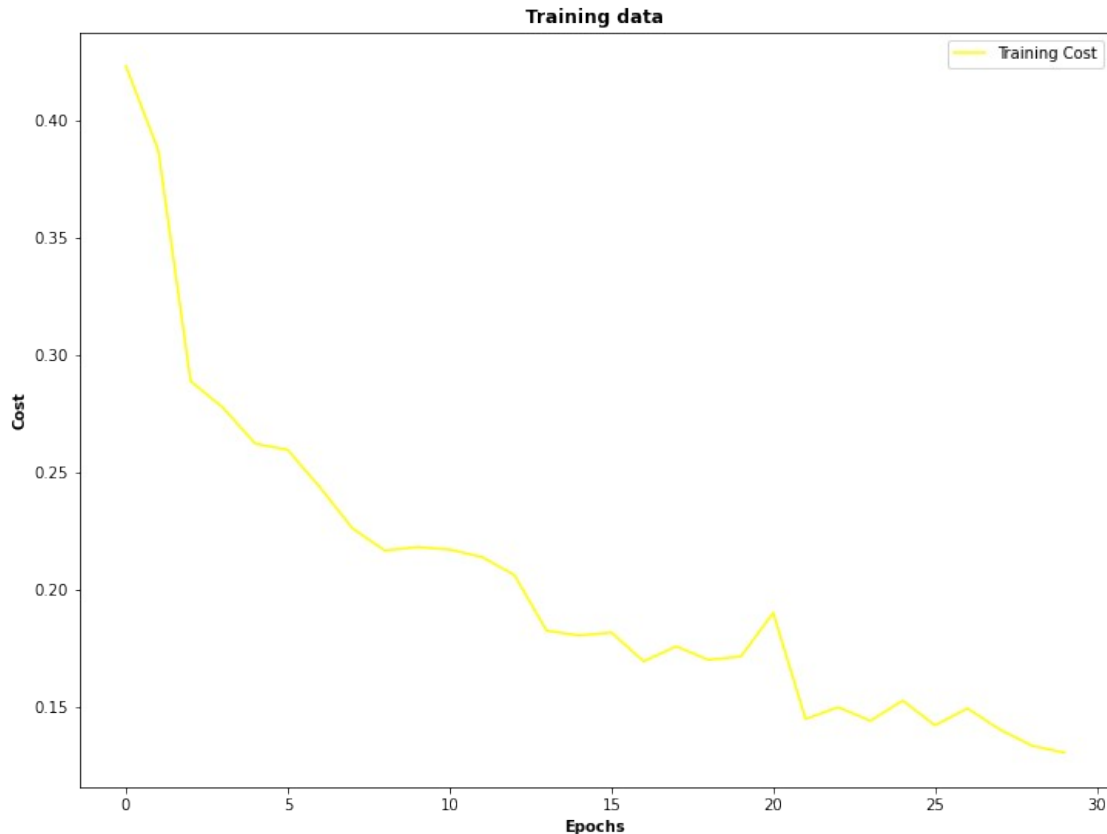
```

monitor_training_cost = True,
monitor_training_accuracy = False,
monitor_evaluation_cost = True,
monitor_evaluation_accuracy = False,

```

```
plot_cost = True,  
evaluation_data = validation_data)
```





For computational purposes we decide to reduce the size of the training and validation set. As we are comparing different learning rates, we don't currently care about their absolute performance but mainly about their comparative performance, so in the following we are going to explore the figures rather than the specific accuracy on the training and validation set. What we obtain from the following figures is that for $\eta = 0.05, 0.5$ the accuracy is almost better both in training and validation set, but in the validation set the oscillations are omnipresent, henceforth the credibility of the model accuracy is not stable.

```
net1 = Network1(sizes = [784,30,10])
results_accuracy_train = []
results_accuracy_valid = []
learning_rates = [0.05,0.5,5]
COLORS = ['#2A6EA6', '#FFCD33', '#FF7033']
n_epochs = 30
for eta in learning_rates:
    results_accuracy_train.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = eta, lmbd = 0,
    monitor_evaluation_cost = False,
    monitor_training_accuracy = True,
    monitor_training_cost = False,
    hyper_training_accuracy_plot = True,
    hyper_evaluation_accuracy_plot =
False,
```



```

                                evaluation_data =
validation_data[:500]))
    results_accuracy_valid.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = eta, lmbd = 0,
                                monitor_evaluation_accuracy = True,
                                monitor_training_cost = False,
                                hyper_training_cost_plot = False,
                                hyper_evaluation_accuracy_plot =
True,
                                evaluation_data =
validation_data[:500]))

fig,ax = plt.subplots(figsize = (12,9))
fig1,ax1 = plt.subplots(figsize = (12,9))
for result, eta, color in zip(results_accuracy_train, learning_rates,
COLORS):
    ax.plot(np.arange(n_epochs),[x/1000 for x in result],'o-',label
= "$\eta$ = "+str(eta),color = color)
ax.set_xticks(ticks = list(range(n_epochs)))
ax.set_xticklabels(list(range(n_epochs)))
ax.set_xlabel('Epochs', weight = 'bold')
ax.set_ylabel('Accuracy', weight = 'bold')
ax.set_title('Training data', weight = 'bold')
ax.legend(loc = 0)

for result, eta, color in zip(results_accuracy_valid, learning_rates,
COLORS):
    ax1.plot(np.arange(n_epochs),[x/500 for x in result],'o-',label
= "$\eta$ = "+str(eta),color = color)
ax1.set_xticks(ticks = list(range(n_epochs)))
ax1.set_xticklabels(list(range(n_epochs)))
ax1.set_xlabel('Epochs', weight = 'bold')
ax1.set_ylabel('Accuracy', weight = 'bold')
ax1.set_title('Validation data', weight = 'bold')
ax1.legend(loc = 0)
plt.show()

```



```


# Let us now generate the respective figures for the cost function.
# Here we see that  $\eta = 0.5$  outperforms the other values on
# the training data, while  $\eta = 5$  comes second but with many
# fluctuations indicating that this value may be too large. On the
# other side, on validation set all figures have increasing trend
# indicating that the model does not perform in the desired way
# while trying to classify the 'unseen' data.
net1 = Network1(sizes = [784,30,10])
results_cost_train = []
results_cost_valid = []
learning_rates = [0.05,0.5,5]
COLORS = ['#2A6EA6', '#FFCD33', '#FF7033']
n_epochs = 30
for eta in learning_rates:
    results_cost_train.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = eta, lmbd = 0,
                                monitor_evaluation_cost = False,
                                monitor_training_cost = True,
                                hyper_training_cost_plot = True,
                                hyper_evaluation_cost_plot = False,
                                evaluation_data =
validation_data[:500]))
    results_cost_valid.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = eta, lmbd = 0,
                                monitor_evaluation_cost = True,
                                monitor_training_cost = False,
                                hyper_training_cost_plot = False,
                                hyper_evaluation_cost_plot = True,
                                evaluation_data =
validation_data[:500]))

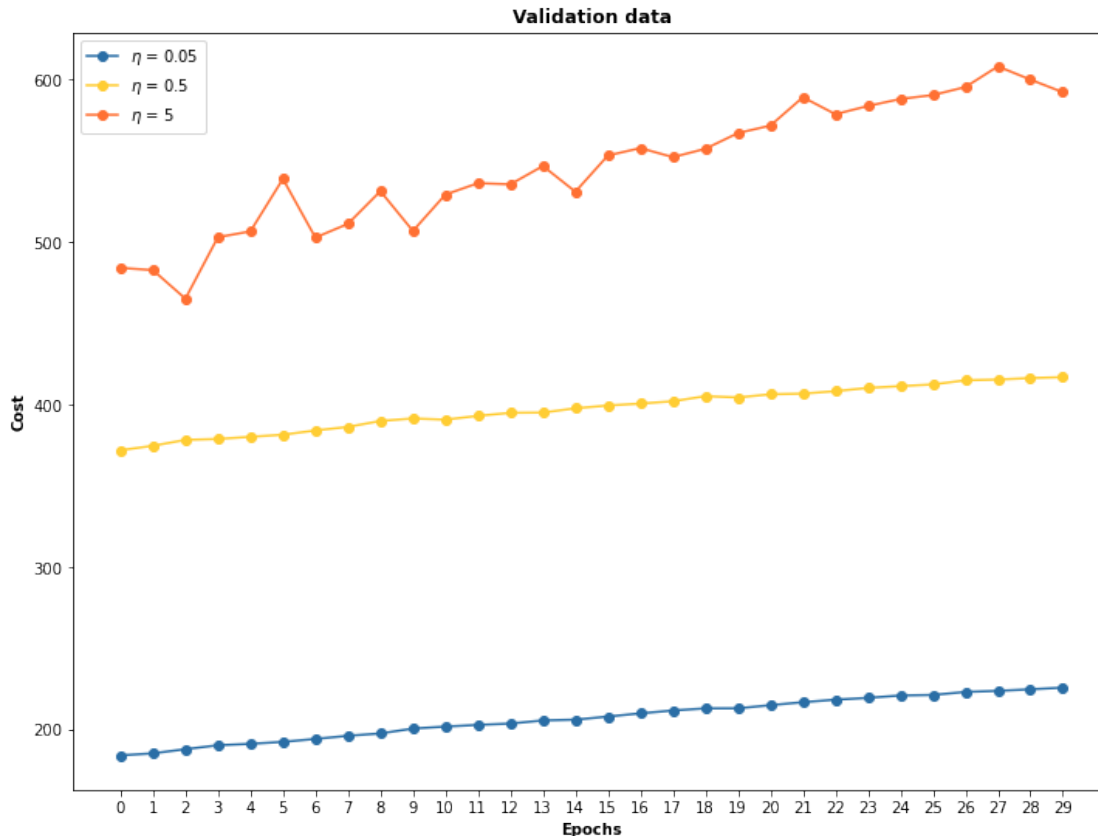
fig,ax = plt.subplots(figsize = (12,9))
fig1,ax1 = plt.subplots(figsize = (12,9))
for result, eta, color in zip(results_cost_train, learning_rates,
COLORS):
    ax.plot(np.arange(n_epochs),result,'o-',label = "$\eta$ =
"+str(eta),color = color)
ax.set_xticks(ticks = list(range(n_epochs)))
ax.set_xticklabels(list(range(n_epochs)))
ax.set_xlabel('Epochs', weight = 'bold')
ax.set_ylabel('Cost', weight = 'bold')
ax.set_title('Training data', weight = 'bold')
ax.legend(loc = 0)

for result, eta, color in zip(results_cost_valid, learning_rates,
COLORS):
    ax1.plot(np.arange(n_epochs),result,'o-',label = "$\eta$ =
"+str(eta),color = color)
ax1.set_xticks(ticks = list(range(n_epochs)))
ax1.set_xticklabels(list(range(n_epochs)))

```

```
ax1.set_xlabel('Epochs', weight = 'bold')
ax1.set_ylabel('Cost', weight = 'bold')
ax1.set_title('Validation data', weight = 'bold')
ax1.legend(loc = 0)
plt.show()
```





As in three out of four previous graphs learning rate $\eta = 0.5$ has the most desirable outcome, we decide to move forward with this value and now investigate different values for the L2 Regularization term. Firstly, we aim to compare the performance of the model when it is regularised and when its not. Heuristically, we start by comparing $\lambda = 1$ and $\lambda = 0.1$ with no regularization (i.e $\lambda = 0$). So, from the plots below we get that $\lambda = 0$ and $\lambda = 1$ have (almost) identical high accuracy on the training data, while $\lambda = 0.1$ reaches this level of accuracy after epoch 18. Now in terms of the validation set, for all different λ value there are plenty of oscilations, with $\lambda = 0.1$ providing marginally the highest accuracy in time.

```
net1 = Network1(sizes = [784,30,10])
results_accuracy_train = []
results_accuracy_valid = []
regul_param = [0.1,1,0]
COLORS = ['#2A6EA6', '#FFCD33', '#FF7033']
n_epochs = 30
for lmbd in regul_param:
    results_accuracy_train.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = 0.5, lmbd = lmbd,
monitor_evaluation_accuracy = False,
monitor_training_accuracy = True,
hyper_training_accuracy_plot = True,
```


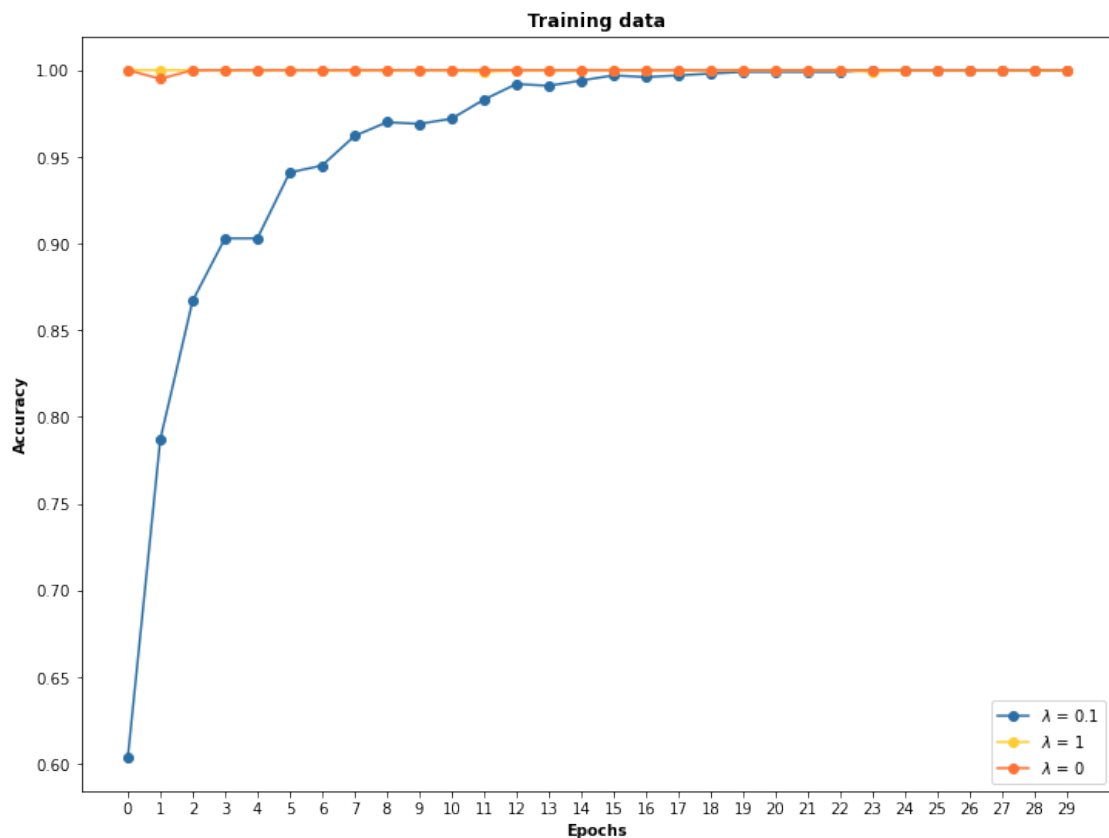
```

hyper_evaluation_accuracy_plot =
False,
evaluation_data =
validation_data[:500]))
    results_accuracy_valid.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = 0.5, lmbd = lmbd,
    monitor_evaluation_accuracy = True,
    monitor_training_accuracy = False,
    hyper_training_accuracy_plot = False,
    hyper_evaluation_accuracy_plot =
True,
evaluation_data =
validation_data[:500]))

fig,ax = plt.subplots(figsize = (12,9))
fig1,ax1 = plt.subplots(figsize = (12,9))
for result, lmbd, color in zip(results_accuracy_train, regul_param,
COLORS):
    ax.plot(np.arange(n_epochs),[x/1000 for x in result],'o-',label
= "$\lambda$ = "+str(lmbd),color = color)
ax.set_xticks(ticks = list(range(n_epochs)))
ax.set_xticklabels(list(range(n_epochs)))
ax.set_xlabel('Epochs', weight = 'bold')
ax.set_ylabel('Accuracy', weight = 'bold')
ax.set_title('Training data', weight = 'bold')
ax.legend(loc = 0)

for result, lmbd, color in zip(results_accuracy_valid, regul_param,
COLORS):
    ax1.plot(np.arange(n_epochs),[x/500 for x in result],'o-',label
= "$\lambda$ = "+str(lmbd),color = color)
ax1.set_xticks(ticks = list(range(n_epochs)))
ax1.set_xticklabels(list(range(n_epochs)))
ax1.set_xlabel('Epochs', weight = 'bold')
ax1.set_ylabel('Accuracy', weight = 'bold')
ax1.set_title('Validation data', weight = 'bold')
ax1.legend(loc = 0)
plt.show()

```




```

# Lets make the cost plots now.
# We see that the highest value of  $\lambda$  (i.e 1) yields the highest cost on
# training data and the lowest on the validation data
# and after trying other triples too, the same pattern hold for all of
# the highest values. This may be the case because greater  $\lambda$ 
# values try to keep weights lower than other values, which in turn
# may lead to better and more controlable weight modifications.
# So, as the weights are getting updated based on the training data,
# when there is less restriction in their magnitude (i.e lower
#  $\lambda$  values) they can be more well fitted to the seen data, but an
# opposite pattern holds for the 'unseen' data. However, as in the
# accuracy graphs above the largest  $\lambda$  value was outperformed in both
# the training and validation set we prefer to continue with
# a small value and in particular with  $\lambda = 0.1$ . However, the
# discrepancy between  $\lambda = 0.1$  and no regularization at all is small, so
# we continue with these two values and we compare them with values of
# less magnitude. Note that since we are using a small part
# of the training and validation dataset the  $\lambda$  value may have to be
# adjusted accordingly if we increase the number of input values.
net1 = Network1(sizes = [784,30,10])
results_cost_train = []
results_cost_valid = []
regul_param = [0.1,1,0]
COLORS = ['#2A6EA6', '#FFCD33', '#FF7033']
n_epochs = 30
for lmbd in regul_param:
    results_cost_train.append(net1.SGD(training_data[:1000],
    mini_batch_size = 15, epochs = n_epochs, eta = 0.5, lmbd = lmbd,
    monitor_evaluation_cost = False,
    monitor_training_cost = True,
    hyper_training_cost_plot = True,
    hyper_evaluation_cost_plot = False,
    evaluation_data =
validation_data[:500]))
    results_cost_valid.append(net1.SGD(training_data[:1000],
    mini_batch_size = 15, epochs = n_epochs, eta = 0.5, lmbd = lmbd,
    monitor_evaluation_cost = True,
    monitor_training_cost = False,
    hyper_training_cost_plot = False,
    hyper_evaluation_cost_plot = True,
    evaluation_data =
validation_data[:500]))

fig,ax = plt.subplots(figsize = (12,9))
fig1,ax1 = plt.subplots(figsize = (12,9))
for result, lmbd, color in zip(results_cost_train, regul_param,
COLORS):
    ax.plot(np.arange(n_epochs),result,'o-',label = "$\lambda$ =
"+str(lmbd),color = color)
ax.set_xticks(ticks = list(range(n_epochs)))

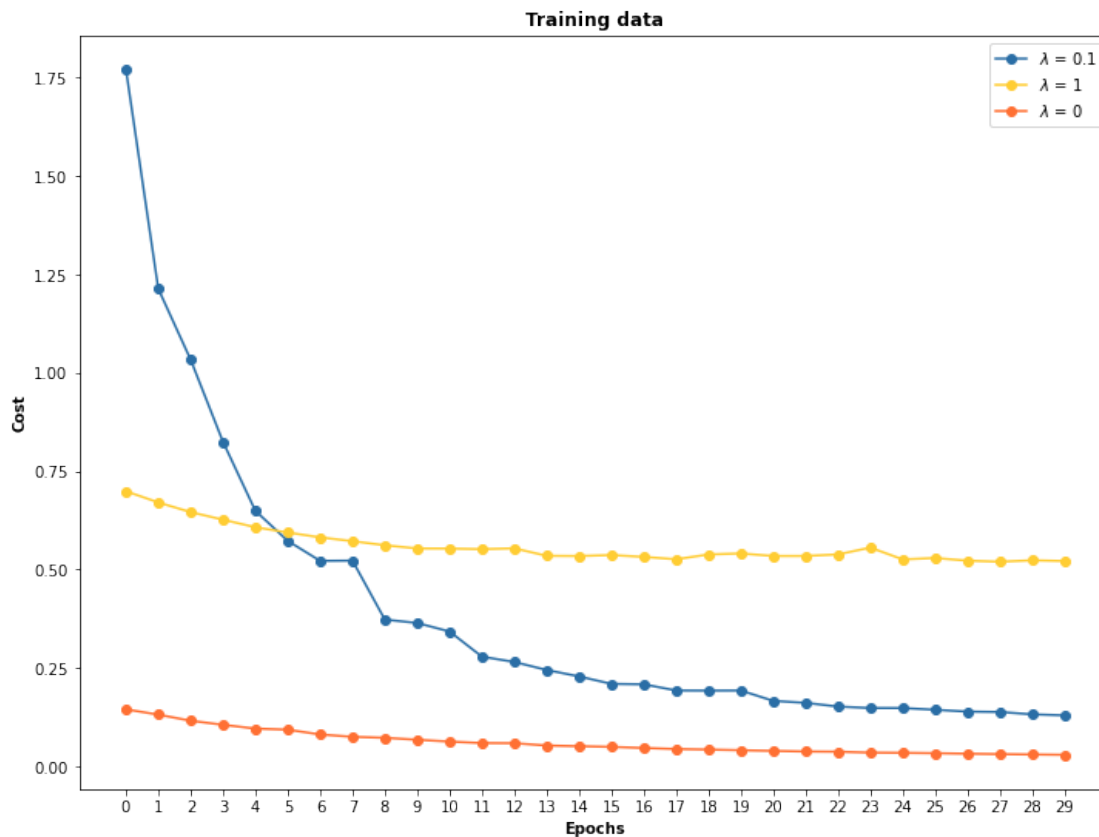
```

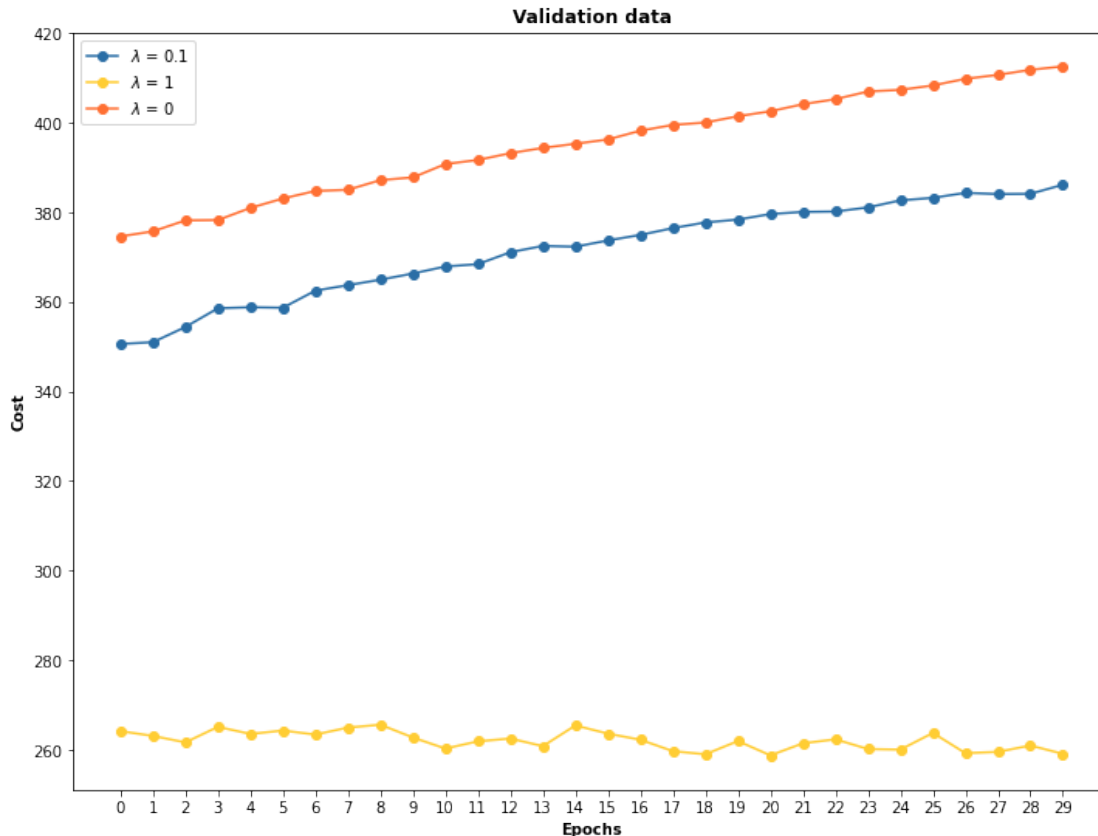
```

ax.set_xticklabels(list(range(n_epochs)))
ax.set_xlabel('Epochs', weight = 'bold')
ax.set_ylabel('Cost', weight = 'bold')
ax.set_title('Training data', weight = 'bold')
ax.legend(loc = 0)

for result, lmbd, color in zip(results_cost_valid, regul_param,
COLORS):
    ax1.plot(np.arange(n_epochs),result,'o-',label = "$\lambda$ = "
"+str(lmbd),color = color)
ax1.set_xticks(ticks = list(range(n_epochs)))
ax1.set_xticklabels(list(range(n_epochs)))
ax1.set_xlabel('Epochs', weight = 'bold')
ax1.set_ylabel('Cost', weight = 'bold')
ax1.set_title('Validation data', weight = 'bold')
ax1.legend(loc = 0)
plt.show()

```





The new lambda values, i.e $\lambda = 0.01$ has equal high accuracy on the training data as $\lambda = 1$, while for $\lambda = 0.1$ the model is improving along epochs. On the other hand, the accuracy on the validation set shows again many fluctuations with the value $\lambda = 1$ reaching the highest and the minimum levels and the $\lambda = 0.01$ having the less variant and with a slight increase performance.

```
net1 = Network1(sizes = [784,30,10])
results_accuracy_train = []
results_accuracy_valid = []
regul_param = [0.1,1,0.01]
COLORS = ['#2A6EA6', '#FFCD33', '#FF7033']
n_epochs = 30
for lmbd in regul_param:
    results_accuracy_train.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = 0.5, lmbd = lmbd,
monitor_evaluation_accuracy = False,
monitor_training_accuracy = True,
hyper_training_accuracy_plot = True,
hyper_evaluation_accuracy_plot =
False,
evaluation_data =
validation_data[:500]))
    results_accuracy_valid.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = 0.5, lmbd = lmbd,
```


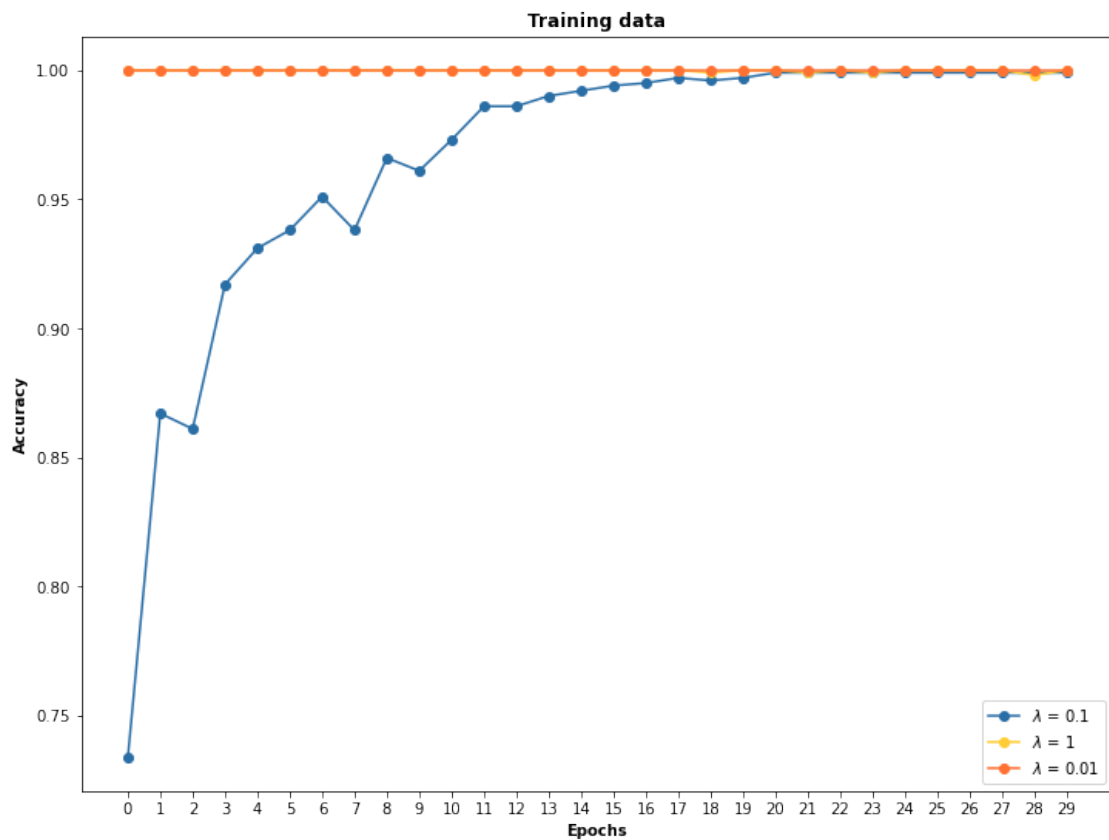
```

monitor_evaluation_accuracy = True,
monitor_training_accuracy = False,
hyper_training_accuracy_plot = False,
hyper_evaluation_accuracy_plot =
True,
evaluation_data =
validation_data[:500]))

fig,ax = plt.subplots(figsize = (12,9))
fig1,ax1 = plt.subplots(figsize = (12,9))
for result, lmbd, color in zip(results_accuracy_train, regul_param,
COLORS):
    ax.plot(np.arange(n_epochs),[x/1000 for x in result],'o-',label
= "$\lambda$ = "+str(lmbd),color = color)
ax.set_xticks(ticks = list(range(n_epochs)))
ax.set_xticklabels(list(range(n_epochs)))
ax.set_xlabel('Epochs', weight = 'bold')
ax.set_ylabel('Accuracy', weight = 'bold')
ax.set_title('Training data', weight = 'bold')
ax.legend(loc = 0)

for result, lmbd, color in zip(results_accuracy_valid, regul_param,
COLORS):
    ax1.plot(np.arange(n_epochs),[x/500 for x in result],'o-',label
= "$\lambda$ = "+str(lmbd),color = color)
ax1.set_xticks(ticks = list(range(n_epochs)))
ax1.set_xticklabels(list(range(n_epochs)))
ax1.set_xlabel('Epochs', weight = 'bold')
ax1.set_ylabel('Accuracy', weight = 'bold')
ax1.set_title('Validation data', weight = 'bold')
ax1.legend(loc = 0)
plt.show()

```




*# As regards the cost of the model, again greater λ values are the best on validation set but the worst on the training set, while
only $\lambda = 0.01$ shows some significant learning process, at least on the training data.*


```
net1 = Network1(sizes = [784,30,10])
results_cost_train = []
results_cost_valid = []
regul_param = [0.1,1,0.01]
COLORS = ['#2A6EA6', '#FFCD33', '#FF7033']
n_epochs = 30
for lmbd in regul_param:
    results_cost_train.append(net1.SGD(training_data[:1000],
    mini_batch_size = 15, epochs = n_epochs, eta = 0.5, lmbd = lmbd,
    monitor_evaluation_cost = False,
    monitor_training_cost = True,
    hyper_training_cost_plot = True,
    hyper_evaluation_cost_plot = False,
    evaluation_data =
validation_data[:500]))
    results_cost_valid.append(net1.SGD(training_data[:1000],
    mini_batch_size = 15, epochs = n_epochs, eta = 0.5, lmbd = lmbd,
    monitor_evaluation_cost = True,
    monitor_training_cost = False,
    hyper_training_cost_plot = False,
    hyper_evaluation_cost_plot = True,
    evaluation_data =
validation_data[:500]))

fig,ax = plt.subplots(figsize = (12,9))
fig1,ax1 = plt.subplots(figsize = (12,9))
for result, lmbd, color in zip(results_cost_train, regul_param,
COLORS):
    ax.plot(np.arange(n_epochs),result,'o-',label = "$\lambda$ =
"+str(lmbd),color = color)
ax.set_xticks(ticks = list(range(n_epochs)))
ax.set_xticklabels(list(range(n_epochs)))
ax.set_xlabel('Epochs', weight = 'bold')
ax.set_ylabel('Cost', weight = 'bold')
ax.set_title('Training data', weight = 'bold')
ax.legend(loc = 0)

for result, lmbd, color in zip(results_cost_valid, regul_param,
COLORS):
    ax1.plot(np.arange(n_epochs),result,'o-',label = "$\lambda$ =
"+str(lmbd),color = color)
ax1.set_xticks(ticks = list(range(n_epochs)))
ax1.set_xticklabels(list(range(n_epochs)))
ax1.set_xlabel('Epochs', weight = 'bold')
ax1.set_ylabel('Cost', weight = 'bold')
ax1.set_title('Validation data', weight = 'bold')
```

```
ax1.legend(loc = 0)  
plt.show()
```





In the following we consider the whole training and validation set, we adjust $\lambda = 0.1$ to this, and we measure the accuracy of this # model on the validation set. What we obtain, is that there is no improvment compared to the accuracy of the model without # L2 regularization.q

```
net = Network1(sizes = [784,30,10], cost = cross_entropy)
net.SGD(training_data, mini_batch_size = 15, epochs = 30, eta = 0.5,
        lmbd = 0.1,
        monitor_training_cost = False,
        monitor_training_accuracy = False,
        monitor_evaluation_cost = False,
        monitor_evaluation_accuracy = True,
        evaluation_data = validation_data)
```

```
Accuracy on evaluation data-Epoch0: 9439 / 10000
Accuracy on evaluation data-Epoch1: 9509 / 10000
Accuracy on evaluation data-Epoch2: 9496 / 10000
Accuracy on evaluation data-Epoch3: 9535 / 10000
Accuracy on evaluation data-Epoch4: 9584 / 10000
Accuracy on evaluation data-Epoch5: 9570 / 10000
Accuracy on evaluation data-Epoch6: 9598 / 10000
Accuracy on evaluation data-Epoch7: 9569 / 10000
Accuracy on evaluation data-Epoch8: 9504 / 10000
Accuracy on evaluation data-Epoch9: 9620 / 10000
Accuracy on evaluation data-Epoch10: 9552 / 10000
```



```
Accuracy on evaluation data-Epoch11: 9600 / 10000
Accuracy on evaluation data-Epoch12: 9609 / 10000
Accuracy on evaluation data-Epoch13: 9613 / 10000
Accuracy on evaluation data-Epoch14: 9566 / 10000
Accuracy on evaluation data-Epoch15: 9598 / 10000
Accuracy on evaluation data-Epoch16: 9569 / 10000
Accuracy on evaluation data-Epoch17: 9599 / 10000
Accuracy on evaluation data-Epoch18: 9584 / 10000
Accuracy on evaluation data-Epoch19: 9589 / 10000
Accuracy on evaluation data-Epoch20: 9593 / 10000
Accuracy on evaluation data-Epoch21: 9598 / 10000
Accuracy on evaluation data-Epoch22: 9572 / 10000
Accuracy on evaluation data-Epoch23: 9584 / 10000
Accuracy on evaluation data-Epoch24: 9593 / 10000
Accuracy on evaluation data-Epoch25: 9582 / 10000
Accuracy on evaluation data-Epoch26: 9591 / 10000
Accuracy on evaluation data-Epoch27: 9565 / 10000
Accuracy on evaluation data-Epoch28: 9589 / 10000
Accuracy on evaluation data-Epoch29: 9596 / 10000
```

*# We continue with $\lambda = 0.1$ but we double the number of epochs as this λ value indicated some accuracy progress throughout epochs
on the training data. However we see that finally there is no improvement in the accuracy and thus other parameters must be found.*

```
net = Network1(sizes = [784,30,10], cost = cross_entropy)
net.SGD(training_data, mini_batch_size = 15, epochs = 60, eta = 0.5,
        lmbd = 0.1,
        monitor_training_cost = False,
        monitor_training_accuracy = False,
        monitor_evaluation_cost = False,
        monitor_evaluation_accuracy = True,
        evaluation_data = validation_data)
```

```
Accuracy on evaluation data-Epoch0: 9401 / 10000
Accuracy on evaluation data-Epoch1: 9475 / 10000
Accuracy on evaluation data-Epoch2: 9537 / 10000
Accuracy on evaluation data-Epoch3: 9545 / 10000
Accuracy on evaluation data-Epoch4: 9532 / 10000
Accuracy on evaluation data-Epoch5: 9613 / 10000
Accuracy on evaluation data-Epoch6: 9585 / 10000
Accuracy on evaluation data-Epoch7: 9592 / 10000
Accuracy on evaluation data-Epoch8: 9601 / 10000
Accuracy on evaluation data-Epoch9: 9606 / 10000
Accuracy on evaluation data-Epoch10: 9576 / 10000
Accuracy on evaluation data-Epoch11: 9605 / 10000
Accuracy on evaluation data-Epoch12: 9564 / 10000
Accuracy on evaluation data-Epoch13: 9602 / 10000
Accuracy on evaluation data-Epoch14: 9600 / 10000
Accuracy on evaluation data-Epoch15: 9599 / 10000
Accuracy on evaluation data-Epoch16: 9585 / 10000
```

Accuracy on evaluation data-Epoch17: 9578 / 10000
Accuracy on evaluation data-Epoch18: 9607 / 10000
Accuracy on evaluation data-Epoch19: 9608 / 10000
Accuracy on evaluation data-Epoch20: 9612 / 10000
Accuracy on evaluation data-Epoch21: 9600 / 10000
Accuracy on evaluation data-Epoch22: 9610 / 10000
Accuracy on evaluation data-Epoch23: 9588 / 10000
Accuracy on evaluation data-Epoch24: 9629 / 10000
Accuracy on evaluation data-Epoch25: 9586 / 10000
Accuracy on evaluation data-Epoch26: 9587 / 10000
Accuracy on evaluation data-Epoch27: 9591 / 10000
Accuracy on evaluation data-Epoch28: 9582 / 10000
Accuracy on evaluation data-Epoch29: 9599 / 10000
Accuracy on evaluation data-Epoch30: 9592 / 10000
Accuracy on evaluation data-Epoch31: 9583 / 10000
Accuracy on evaluation data-Epoch32: 9597 / 10000
Accuracy on evaluation data-Epoch33: 9608 / 10000
Accuracy on evaluation data-Epoch34: 9608 / 10000
Accuracy on evaluation data-Epoch35: 9609 / 10000
Accuracy on evaluation data-Epoch36: 9607 / 10000
Accuracy on evaluation data-Epoch37: 9606 / 10000
Accuracy on evaluation data-Epoch38: 9585 / 10000
Accuracy on evaluation data-Epoch39: 9590 / 10000
Accuracy on evaluation data-Epoch40: 9595 / 10000
Accuracy on evaluation data-Epoch41: 9598 / 10000
Accuracy on evaluation data-Epoch42: 9598 / 10000
Accuracy on evaluation data-Epoch43: 9600 / 10000
Accuracy on evaluation data-Epoch44: 9606 / 10000
Accuracy on evaluation data-Epoch45: 9563 / 10000
Accuracy on evaluation data-Epoch46: 9603 / 10000
Accuracy on evaluation data-Epoch47: 9607 / 10000
Accuracy on evaluation data-Epoch48: 9572 / 10000
Accuracy on evaluation data-Epoch49: 9595 / 10000
Accuracy on evaluation data-Epoch50: 9592 / 10000
Accuracy on evaluation data-Epoch51: 9595 / 10000
Accuracy on evaluation data-Epoch52: 9594 / 10000
Accuracy on evaluation data-Epoch53: 9587 / 10000
Accuracy on evaluation data-Epoch54: 9579 / 10000
Accuracy on evaluation data-Epoch55: 9536 / 10000
Accuracy on evaluation data-Epoch56: 9579 / 10000
Accuracy on evaluation data-Epoch57: 9585 / 10000
Accuracy on evaluation data-Epoch58: 9591 / 10000
Accuracy on evaluation data-Epoch59: 9570 / 10000

Now, instead of trying one of the η and λ each time we decide to monitor their performance simultaneously as follows.

*# We observe that the only combination that showcases some desirable development both on training and validation set through
epochs is the $(\lambda, \eta) = (5, 0.02)$ and it has finally best cost on the*

```

validation set. The combination (0.1,3) has the lowest cost
# on the training data but the highest and most abrupt in terms of
changes on the validation set.
# We observe that lower or medium  $\eta, \lambda$  values lead to higher accuracy
but greater difference between the model performances on
# the training and validation set. For very small learning rates there
is an increasing trend but lower accuracy
# (always up to the number of epochs) What about high  $\lambda$  and low  $\eta$ ?
net1 = Network1(sizes = [784,30,10])
results_cost_train = []
results_cost_valid = []
learn_rates = [0.2,0.5,3,0.02]
regul_param = [2,1,0.1,5]
COLORS = ['#2A6EA6', '#FFCD33', '#FF7033', 'r']
n_epochs = 60
for eta, lmbd in zip(learn_rates, regul_param):
    results_cost_train.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = eta, lmbd = lmbd,
                                monitor_evaluation_accuracy = False,
                                monitor_training_cost = True,
                                hyper_training_cost_plot = True,
                                hyper_evaluation_accuracy_plot =
False,
                                evaluation_data =
validation_data[:500]))
    results_cost_valid.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = eta, lmbd = lmbd,
                                monitor_evaluation_accuracy = False,
                                monitor_evaluation_cost = True,
                                hyper_training_accuracy_plot = False,
                                hyper_evaluation_cost_plot = True,
                                evaluation_data =
validation_data[:500]))

fig, ax = plt.subplots(figsize = (15,9))
fig1, ax1 = plt.subplots(figsize = (15,9))
for result, eta, lmbd, color in zip(results_cost_train,
learn_rates, regul_param, COLORS):
    ax.plot(np.arange(n_epochs), [x for x in result], 'o-', label = "$\lambda$
lambda$, $\eta$ = "+str([lmbd, eta]), color = color)
ax.set_xticks(ticks = list(range(n_epochs)))
ax.set_xticklabels(list(range(n_epochs)))
ax.set_xlabel('Epochs', weight = 'bold')
ax.set_ylabel('Cost', weight = 'bold')
ax.set_title('Training data', weight = 'bold')
ax.legend(loc = 0)


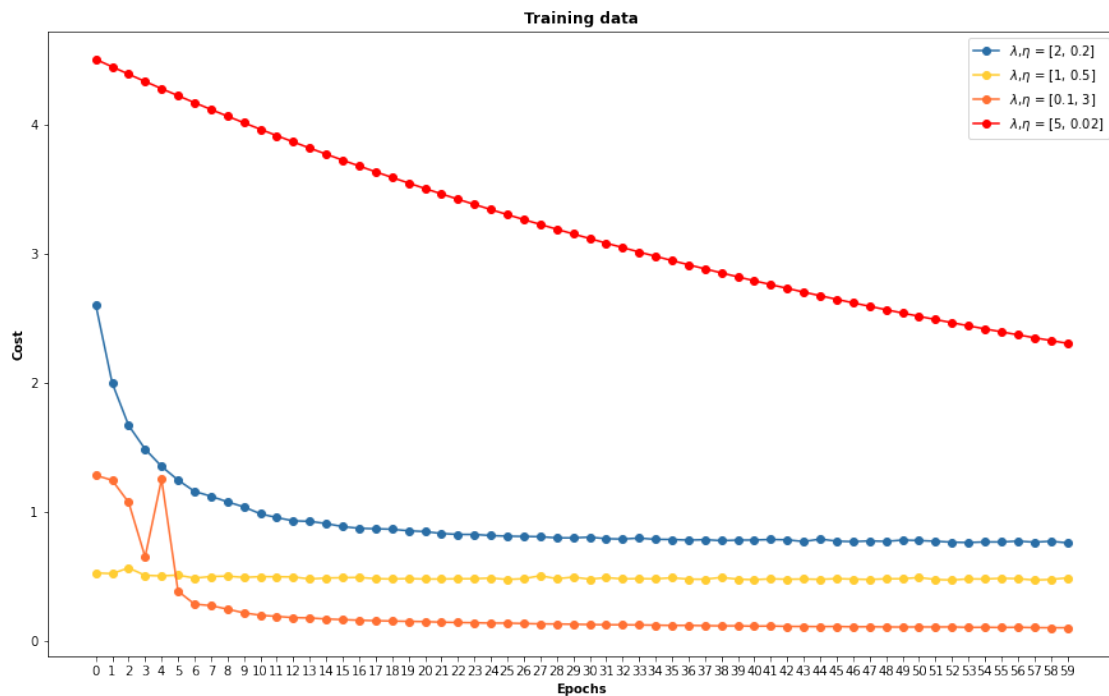
for result, eta, lmbd, color in zip(results_cost_valid, learn_rates,
regul_param, COLORS):
    ax1.plot(np.arange(n_epochs), [x for x in result], 'o-', label =

```

```

"$\lambda$, $\eta$ = "+str([lmbd,eta]),color = color)
ax1.set_xticks(ticks = list(range(n_epochs)))
ax1.set_xticklabels(list(range(n_epochs)))
ax1.set_xlabel('Epochs', weight = 'bold')
ax1.set_ylabel('Cost', weight = 'bold')
ax1.set_title('Validation data', weight = 'bold')
ax1.legend(loc = 0)
plt.show()

```



```

# In the  $\phi\lambda\lambda\sigma\iota\nu\gamma$ , the combination  $(\lambda, \eta) = (2, 0.2)$  demonstrates some
# gradual increase on the training data while the rest of the
# combinations reach their highest accuracy from the  $\omega\epsilon\pi\upsilon$  first
# epochs. On the validation set the combination  $(5, 0.02)$  shows
# the least variant behaviour but with no increasing disposition.
net1 = Network1(sizes = [784, 30, 10])
results_accuracy_train = []
results_accuracy_valid = []
learn_rates = [0.2, 0.5, 3, 0.02]
regul_param = [2, 1, 0.1, 5]
COLORS = ['#2A6EA6', '#FFCD33', '#FF7033', 'r']
n_epochs = 60
for eta, lmbd in zip(learn_rates, regul_param):
    results_accuracy_train.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = eta, lmbd = lmbd,
monitor_evaluation_accuracy = False,
monitor_training_accuracy = True,
hyper_training_accuracy_plot = True,
hyper_evaluation_accuracy_plot =
False,
evaluation_data =
validation_data[:500]))
    results_accuracy_valid.append(net1.SGD(training_data[:1000],
mini_batch_size = 15, epochs = n_epochs, eta = eta, lmbd = lmbd,
monitor_evaluation_accuracy = True,
monitor_training_accuracy = False,
hyper_training_accuracy_plot = False,
hyper_evaluation_accuracy_plot =
True,
evaluation_data =
validation_data[:500]))

fig, ax = plt.subplots(figsize = (15, 9))
fig1, ax1 = plt.subplots(figsize = (15, 9))
for result, eta, lmbd, color in zip(results_accuracy_train,
learn_rates, regul_param, COLORS):
    ax.plot(np.arange(n_epochs), [x/1000 for x in result], 'o-', label
= "$\lambda$, $\eta$ = "+str([lmbd, eta]), color = color)
ax.set_xticks(ticks = list(range(n_epochs)))
ax.set_xticklabels(list(range(n_epochs)))
ax.set_xlabel('Epochs', weight = 'bold')
ax.set_ylabel('Accuracy', weight = 'bold')
ax.set_title('Training data', weight = 'bold')
ax.legend(loc = 0)


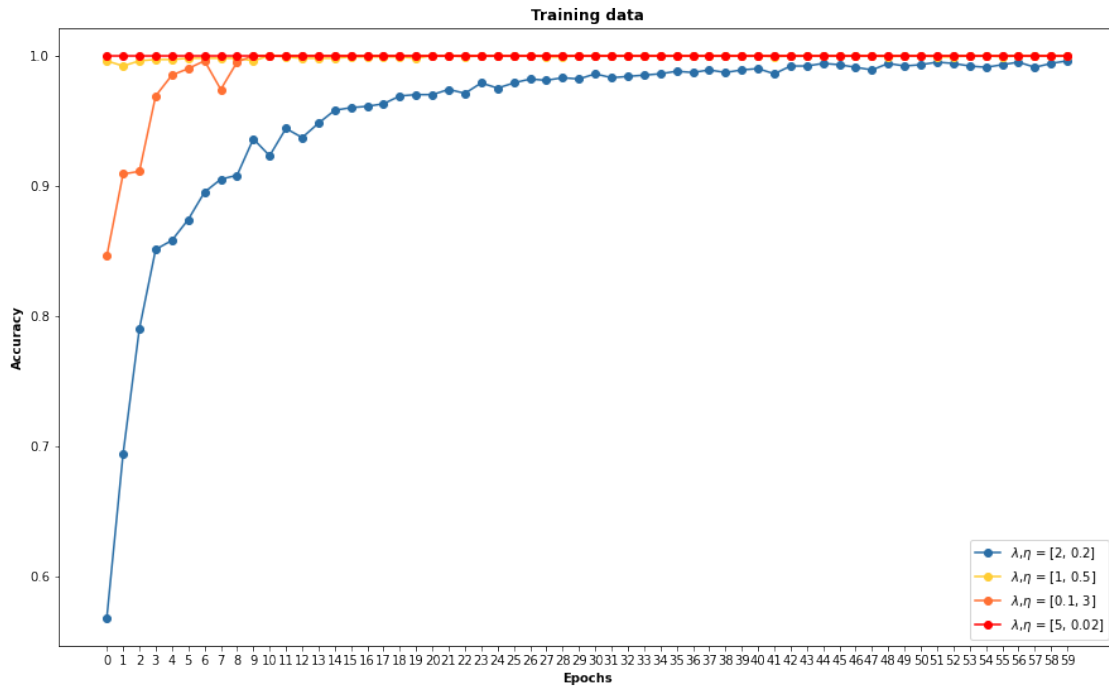
for result, eta, lmbd, color in zip(results_accuracy_valid, learn_rates,
regul_param, COLORS):
    ax1.plot(np.arange(n_epochs), [x/500 for x in result], 'o-', label
= "$\lambda$, $\eta$ = "+str([lmbd, eta]), color = color)
ax1.set_xticks(ticks = list(range(n_epochs)))

```

```

ax1.set_xticklabels(list(range(n_epochs)))
ax1.set_xlabel('Epochs', weight = 'bold')
ax1.set_ylabel('Accuracy', weight = 'bold')
ax1.set_title('Validation data', weight = 'bold')
ax1.legend(loc = 0)
plt.show()

```



```

# From the above the two combinations that seem to learn along the
epochs are the  $(\lambda, \eta) = (5, 0.02)$ ,  $(2, 0.2)$  and therefore
# we proceed to get their accuracy level on the validation set, using
the aggregate values of the sets. The following values
# for  $(5, 0.02)$  have a desirable increasing trend in general but of
small rate and finally there is no improvement compared
# to the previous models(including the no regularization one).
net = Network1(sizes = [784, 30, 10], cost = cross_entropy)
net.SGD(training_data, mini_batch_size = 15, epochs = 60, eta = 0.02,
lmbd = 5,

```

```

    monitor_training_cost = False,
    monitor_training_accuracy = False,
    monitor_evaluation_cost = False,
    monitor_evaluation_accuracy = True,
    evaluation_data = validation_data)

```

```

Accuracy on evaluation data-Epoch0: 8728 / 10000
Accuracy on evaluation data-Epoch1: 9026 / 10000
Accuracy on evaluation data-Epoch2: 9126 / 10000
Accuracy on evaluation data-Epoch3: 9177 / 10000
Accuracy on evaluation data-Epoch4: 9239 / 10000
Accuracy on evaluation data-Epoch5: 9272 / 10000
Accuracy on evaluation data-Epoch6: 9301 / 10000
Accuracy on evaluation data-Epoch7: 9332 / 10000
Accuracy on evaluation data-Epoch8: 9338 / 10000
Accuracy on evaluation data-Epoch9: 9377 / 10000
Accuracy on evaluation data-Epoch10: 9394 / 10000
Accuracy on evaluation data-Epoch11: 9412 / 10000
Accuracy on evaluation data-Epoch12: 9444 / 10000
Accuracy on evaluation data-Epoch13: 9448 / 10000
Accuracy on evaluation data-Epoch14: 9456 / 10000
Accuracy on evaluation data-Epoch15: 9468 / 10000
Accuracy on evaluation data-Epoch16: 9479 / 10000
Accuracy on evaluation data-Epoch17: 9488 / 10000
Accuracy on evaluation data-Epoch18: 9501 / 10000
Accuracy on evaluation data-Epoch19: 9503 / 10000
Accuracy on evaluation data-Epoch20: 9518 / 10000
Accuracy on evaluation data-Epoch21: 9525 / 10000
Accuracy on evaluation data-Epoch22: 9520 / 10000
Accuracy on evaluation data-Epoch23: 9529 / 10000
Accuracy on evaluation data-Epoch24: 9533 / 10000
Accuracy on evaluation data-Epoch25: 9548 / 10000
Accuracy on evaluation data-Epoch26: 9537 / 10000
Accuracy on evaluation data-Epoch27: 9540 / 10000
Accuracy on evaluation data-Epoch28: 9547 / 10000
Accuracy on evaluation data-Epoch29: 9533 / 10000
Accuracy on evaluation data-Epoch30: 9550 / 10000
Accuracy on evaluation data-Epoch31: 9547 / 10000
Accuracy on evaluation data-Epoch32: 9566 / 10000
Accuracy on evaluation data-Epoch33: 9553 / 10000

```



```

Accuracy on evaluation data-Epoch34: 9558 / 10000
Accuracy on evaluation data-Epoch35: 9551 / 10000
Accuracy on evaluation data-Epoch36: 9571 / 10000
Accuracy on evaluation data-Epoch37: 9574 / 10000
Accuracy on evaluation data-Epoch38: 9557 / 10000
Accuracy on evaluation data-Epoch39: 9570 / 10000
Accuracy on evaluation data-Epoch40: 9564 / 10000
Accuracy on evaluation data-Epoch41: 9577 / 10000
Accuracy on evaluation data-Epoch42: 9561 / 10000
Accuracy on evaluation data-Epoch43: 9572 / 10000
Accuracy on evaluation data-Epoch44: 9576 / 10000
Accuracy on evaluation data-Epoch45: 9578 / 10000
Accuracy on evaluation data-Epoch46: 9580 / 10000
Accuracy on evaluation data-Epoch47: 9575 / 10000
Accuracy on evaluation data-Epoch48: 9574 / 10000
Accuracy on evaluation data-Epoch49: 9585 / 10000
Accuracy on evaluation data-Epoch50: 9586 / 10000
Accuracy on evaluation data-Epoch51: 9585 / 10000
Accuracy on evaluation data-Epoch52: 9589 / 10000
Accuracy on evaluation data-Epoch53: 9588 / 10000
Accuracy on evaluation data-Epoch54: 9593 / 10000
Accuracy on evaluation data-Epoch55: 9581 / 10000
Accuracy on evaluation data-Epoch56: 9588 / 10000
Accuracy on evaluation data-Epoch57: 9597 / 10000
Accuracy on evaluation data-Epoch58: 9605 / 10000
Accuracy on evaluation data-Epoch59: 9600 / 10000

```

Here the combination (2,0.2) ameliorates the performance of the model reaching the rate of 9688/10000 which is better than the 9667/10000 of the model without regularization. However is not keeping improving towards the last epoch.

Even if we have doubled the number of epochs, this model outperforms the no regularization model within the first 30 epochs as well. So we decide to continue with this model in the following step.

```

net = Network1(sizes = [784,30,10], cost = cross_entropy)
net.SGD(training_data, mini_batch_size = 15, epochs = 60, eta = 0.2,
        lmbd = 2,
        monitor_training_cost = False,
        monitor_training_accuracy = False,
        monitor_evaluation_cost = False,
        monitor_evaluation_accuracy = True,
        evaluation_data = validation_data)

```

```

Accuracy on evaluation data-Epoch0: 9311 / 10000
Accuracy on evaluation data-Epoch1: 9474 / 10000
Accuracy on evaluation data-Epoch2: 9514 / 10000
Accuracy on evaluation data-Epoch3: 9494 / 10000
Accuracy on evaluation data-Epoch4: 9568 / 10000
Accuracy on evaluation data-Epoch5: 9605 / 10000
Accuracy on evaluation data-Epoch6: 9607 / 10000

```


Accuracy on evaluation data-Epoch7: 9605 / 10000
Accuracy on evaluation data-Epoch8: 9632 / 10000
Accuracy on evaluation data-Epoch9: 9612 / 10000
Accuracy on evaluation data-Epoch10: 9640 / 10000
Accuracy on evaluation data-Epoch11: 9629 / 10000
Accuracy on evaluation data-Epoch12: 9655 / 10000
Accuracy on evaluation data-Epoch13: 9651 / 10000
Accuracy on evaluation data-Epoch14: 9652 / 10000
Accuracy on evaluation data-Epoch15: 9647 / 10000
Accuracy on evaluation data-Epoch16: 9665 / 10000
Accuracy on evaluation data-Epoch17: 9637 / 10000
Accuracy on evaluation data-Epoch18: 9662 / 10000
Accuracy on evaluation data-Epoch19: 9649 / 10000
Accuracy on evaluation data-Epoch20: 9661 / 10000
Accuracy on evaluation data-Epoch21: 9665 / 10000
Accuracy on evaluation data-Epoch22: 9648 / 10000
Accuracy on evaluation data-Epoch23: 9668 / 10000
Accuracy on evaluation data-Epoch24: 9660 / 10000
Accuracy on evaluation data-Epoch25: 9672 / 10000
Accuracy on evaluation data-Epoch26: 9680 / 10000
Accuracy on evaluation data-Epoch27: 9685 / 10000
Accuracy on evaluation data-Epoch28: 9652 / 10000
Accuracy on evaluation data-Epoch29: 9661 / 10000
Accuracy on evaluation data-Epoch30: 9670 / 10000
Accuracy on evaluation data-Epoch31: 9657 / 10000
Accuracy on evaluation data-Epoch32: 9675 / 10000
Accuracy on evaluation data-Epoch33: 9631 / 10000
Accuracy on evaluation data-Epoch34: 9664 / 10000
Accuracy on evaluation data-Epoch35: 9666 / 10000
Accuracy on evaluation data-Epoch36: 9658 / 10000
Accuracy on evaluation data-Epoch37: 9688 / 10000
Accuracy on evaluation data-Epoch38: 9680 / 10000
Accuracy on evaluation data-Epoch39: 9666 / 10000
Accuracy on evaluation data-Epoch40: 9650 / 10000
Accuracy on evaluation data-Epoch41: 9638 / 10000
Accuracy on evaluation data-Epoch42: 9682 / 10000
Accuracy on evaluation data-Epoch43: 9665 / 10000
Accuracy on evaluation data-Epoch44: 9671 / 10000
Accuracy on evaluation data-Epoch45: 9674 / 10000
Accuracy on evaluation data-Epoch46: 9656 / 10000
Accuracy on evaluation data-Epoch47: 9641 / 10000
Accuracy on evaluation data-Epoch48: 9658 / 10000
Accuracy on evaluation data-Epoch49: 9680 / 10000
Accuracy on evaluation data-Epoch50: 9660 / 10000
Accuracy on evaluation data-Epoch51: 9671 / 10000
Accuracy on evaluation data-Epoch52: 9668 / 10000
Accuracy on evaluation data-Epoch53: 9657 / 10000
Accuracy on evaluation data-Epoch54: 9674 / 10000
Accuracy on evaluation data-Epoch55: 9680 / 10000
Accuracy on evaluation data-Epoch56: 9668 / 10000

Accuracy on evaluation data-Epoch57: 9671 / 10000
Accuracy on evaluation data-Epoch58: 9673 / 10000
Accuracy on evaluation data-Epoch59: 9659 / 10000

Let us now check this last model to the test data that we were keeping hidden throughout the whole process.

It reaches accuracy level of 96,5% and it has a slight rising tendency.

Here we end the investigation of the models parameters. At this stage I want to mention that L1 regularization techniques were used as well, but without any significant success and therefore the details were omitted. We could continue the checking by incorporating the other parameters too(number of layers, neurons per layer, number of epochs, mini batch size), but the scope of this project was to build and improve by some standard techniques a neural network.

```
net = Network1(sizes = [784,30,10], cost = cross_entropy)
net.SGD(training_data, mini_batch_size = 15, epochs = 60, eta = 0.2,
        lmbd = 2,
        monitor_training_cost = False,
        monitor_training_accuracy = False,
        monitor_evaluation_cost = False,
        monitor_evaluation_accuracy = True,
        evaluation_data = test_data)
```

Accuracy on evaluation data-Epoch0: 9260 / 10000
Accuracy on evaluation data-Epoch1: 9393 / 10000
Accuracy on evaluation data-Epoch2: 9401 / 10000
Accuracy on evaluation data-Epoch3: 9389 / 10000
Accuracy on evaluation data-Epoch4: 9534 / 10000
Accuracy on evaluation data-Epoch5: 9524 / 10000
Accuracy on evaluation data-Epoch6: 9533 / 10000
Accuracy on evaluation data-Epoch7: 9576 / 10000
Accuracy on evaluation data-Epoch8: 9567 / 10000
Accuracy on evaluation data-Epoch9: 9596 / 10000
Accuracy on evaluation data-Epoch10: 9574 / 10000
Accuracy on evaluation data-Epoch11: 9598 / 10000
Accuracy on evaluation data-Epoch12: 9587 / 10000
Accuracy on evaluation data-Epoch13: 9606 / 10000
Accuracy on evaluation data-Epoch14: 9619 / 10000
Accuracy on evaluation data-Epoch15: 9630 / 10000
Accuracy on evaluation data-Epoch16: 9619 / 10000
Accuracy on evaluation data-Epoch17: 9603 / 10000
Accuracy on evaluation data-Epoch18: 9621 / 10000
Accuracy on evaluation data-Epoch19: 9624 / 10000
Accuracy on evaluation data-Epoch20: 9602 / 10000
Accuracy on evaluation data-Epoch21: 9616 / 10000
Accuracy on evaluation data-Epoch22: 9617 / 10000
Accuracy on evaluation data-Epoch23: 9634 / 10000

Accuracy on evaluation data-Epoch24:	9639	/	10000
Accuracy on evaluation data-Epoch25:	9614	/	10000
Accuracy on evaluation data-Epoch26:	9625	/	10000
Accuracy on evaluation data-Epoch27:	9639	/	10000
Accuracy on evaluation data-Epoch28:	9618	/	10000
Accuracy on evaluation data-Epoch29:	9640	/	10000
Accuracy on evaluation data-Epoch30:	9595	/	10000
Accuracy on evaluation data-Epoch31:	9639	/	10000
Accuracy on evaluation data-Epoch32:	9624	/	10000
Accuracy on evaluation data-Epoch33:	9633	/	10000
Accuracy on evaluation data-Epoch34:	9622	/	10000
Accuracy on evaluation data-Epoch35:	9642	/	10000
Accuracy on evaluation data-Epoch36:	9637	/	10000
Accuracy on evaluation data-Epoch37:	9636	/	10000
Accuracy on evaluation data-Epoch38:	9644	/	10000
Accuracy on evaluation data-Epoch39:	9636	/	10000
Accuracy on evaluation data-Epoch40:	9636	/	10000
Accuracy on evaluation data-Epoch41:	9626	/	10000
Accuracy on evaluation data-Epoch42:	9644	/	10000
Accuracy on evaluation data-Epoch43:	9648	/	10000
Accuracy on evaluation data-Epoch44:	9567	/	10000
Accuracy on evaluation data-Epoch45:	9642	/	10000
Accuracy on evaluation data-Epoch46:	9630	/	10000
Accuracy on evaluation data-Epoch47:	9628	/	10000
Accuracy on evaluation data-Epoch48:	9654	/	10000
Accuracy on evaluation data-Epoch49:	9624	/	10000
Accuracy on evaluation data-Epoch50:	9599	/	10000
Accuracy on evaluation data-Epoch51:	9635	/	10000
Accuracy on evaluation data-Epoch52:	9639	/	10000
Accuracy on evaluation data-Epoch53:	9636	/	10000
Accuracy on evaluation data-Epoch54:	9651	/	10000
Accuracy on evaluation data-Epoch55:	9631	/	10000
Accuracy on evaluation data-Epoch56:	9645	/	10000
Accuracy on evaluation data-Epoch57:	9610	/	10000
Accuracy on evaluation data-Epoch58:	9644	/	10000
Accuracy on evaluation data-Epoch59:	9642	/	10000

References : Neural Network and Deep Learning - Michael Nielsen